

Why Are Logarithmic Returns Important in Finance?

Logarithmic returns are important in finance because they provide a more accurate measure of the percentage change in the value of an asset over a period of time. This is particularly important when analyzing financial data because the compounding effect of returns over time can have a significant impact on the value of an asset.

Logarithmic returns are also useful because they are additive. That is, the logarithmic return of a portfolio composed of multiple assets is simply the sum of the logarithmic returns of each individual asset. This makes it easy to calculate the overall performance of a portfolio over a period of time.

Another advantage of using logarithmic returns is that they are normally distributed, which makes them easier to work with mathematically. This is because the natural logarithm of a number is a continuous and smooth function, which means that the resulting distribution of logarithmic returns is also continuous and smooth.

By using logarithmic returns, we can get a more accurate measure of the percentage change in the value of an asset over a period of time, which is crucial when working with financial data.

```
!pip3 install datetime
!pip3 install yfinance
!pip3 install bs4
!pip3 install os
!pip3 install datetime
!pip3 install pandas
!pip3 install pandas_datareader
!pip3 install statsmodels
```

```
import bs4 as bs
import requests
import yfinance as yf
import datetime
```

Below we downloaded all historical data about S&P500 index from 01-05-1981 to our present days/nowadays

```
resp = requests.get('http://en.wikipedia.org/wiki/List_of_S%26P_500_companies')
soup = bs.BeautifulSoup(resp.text, 'lxml')
table = soup.find('table', {'class': 'wikitable sortable'})
tickers = []
for row in table.findAll('tr')[1:]:
    ticker = row.findAll('td')[0].text
    tickers.append(ticker)

tickers = [s.replace('\n', '') for s in tickers]
start = datetime.datetime(1981,5,1)
end = datetime.datetime.now()
data = yf.download(tickers, start=start, end=end)
print(data)
```

[*****100%*****] 503 of 503 completed

2 Failed downloads:
['BRK.B']: Exception('%ticker%: No timezone found, symbol may be delisted')
['BF.B']: Exception('%ticker%: No price data found, symbol may be delisted (1d 1981-05-01 00:0

Price	Adj Close					
Ticker	A	AAL	AAPL	ABBV	ABNB	ABT
Date						
1981-05-01	NaN	NaN	0.097898	NaN	NaN	0.320195 \
1981-05-04	NaN	NaN	0.097467	NaN	NaN	0.320195
1981-05-05	NaN	NaN	0.097036	NaN	NaN	0.315301
1981-05-06	NaN	NaN	0.094448	NaN	NaN	0.313204
1981-05-07	NaN	NaN	0.095742	NaN	NaN	0.317399
...
2024-04-24	137.490005	13.92	169.020004	167.800003	162.839996	106.889999
2024-04-25	136.369995	14.13	169.889999	167.289993	163.009995	106.860001
2024-04-26	137.740005	13.88	169.300003	159.619995	164.229996	107.529999
2024-04-29	139.589996	13.98	173.500000	161.520004	162.250000	107.269997
2024-04-30	137.039993	13.51	170.330002	162.639999	158.570007	105.970001

Price	...					Volume
Ticker	ACGL	ACN	ADBE	ADI	...	WTW
Date						
1981-05-01	NaN	NaN	NaN	0.808018	...	NaN \
1981-05-04	NaN	NaN	NaN	0.800135	...	NaN
1981-05-05	NaN	NaN	NaN	0.796193	...	NaN
1981-05-06	NaN	NaN	NaN	0.804076	...	NaN
1981-05-07	NaN	NaN	NaN	0.811960	...	NaN
...
2024-04-24	93.190002	313.540009	477.119995	196.500000	...	480500.0
2024-04-25	93.050003	309.000000	473.440002	197.940002	...	1476100.0
2024-04-26	90.900002	308.010010	477.559998	201.970001	...	700700.0
2024-04-29	91.250000	303.160004	473.070007	203.899994	...	474300.0
2024-04-30	93.540001	300.910004	462.829987	200.610001	...	486923.0

Price						
Ticker	WY	WYNN	XEL	XOM	XYL	YUM
Date						
1981-05-01	102450	NaN	49600	5251200	NaN	NaN \
1981-05-04	136800	NaN	472400	4275200	NaN	NaN
1981-05-05	87600	NaN	264000	6878400	NaN	NaN
1981-05-06	128700	NaN	67600	7416000	NaN	NaN
1981-05-07	64950	NaN	96400	9086400	NaN	NaN
...
2024-04-24	3019100	1256700.0	4614400	12101200	1053000.0	1909500.0

2024-04-25	2473900	948900.0	6717500	16041000	963600.0	1693100.0
2024-04-26	3616700	1379700.0	6334100	26718400	1112300.0	1979100.0
2024-04-29	2837700	1588300.0	3884600	17309900	1277800.0	2221000.0
2024-04-30	4608360	2625609.0	6804275	20770037	1282551.0	3504826.0

Price

Ticker	ZBH	ZBRA	ZTS
--------	-----	------	-----

Date

1981-05-01	NaN	NaN	NaN
1981-05-04	NaN	NaN	NaN
1981-05-05	NaN	NaN	NaN

```
import pandas as pd
import numpy as np
import yfinance as yf
import matplotlib.pyplot as plt
from pandas_datareader import data
import plotly.express as px
import datetime as dt
import seaborn as sns
```

Now, let’s extract the historical prices Apple company from yahoo finance using a python library yfinance from 01-05-1981 to the present days/nowadays. (AAPL is one of the 500 S&P 500 stocks)

```
sp500_aapl = yf.download('AAPL', start='1981-05-01', end =dt.datetime.now())
sp500_aapl.head()
```

[*****100%*****] 1 of 1 completed

	Open	High	Low	Close	Adj Close	Volume
Date						
1981-05-01	0.126674	0.127790	0.126674	0.126674	0.097898	16553600
1981-05-04	0.126674	0.126674	0.126116	0.126116	0.097467	14448000
1981-05-05	0.126116	0.126116	0.125558	0.125558	0.097036	17539200
1981-05-06	0.122768	0.122768	0.122210	0.122210	0.094448	18950400
1981-05-07	0.123884	0.124442	0.123884	0.123884	0.095742	9363200

```
sp500_aapl.tail()
```

	Open	High	Low	Close	Adj Close	Volume
Date						
2024-04-23	165.350006	167.050003	164.919998	166.899994	166.899994	49537800
2024-04-24	166.539993	169.300003	166.210007	169.020004	169.020004	48251800
2024-04-25	169.529999	170.610001	168.149994	169.889999	169.889999	50558300
2024-04-26	169.880005	171.339996	169.179993	169.300003	169.300003	44525100
2024-04-29	173.369995	176.029999	173.100006	173.500000	173.500000	66162100


sp500_aapl

	Open	High	Low	Close	Adj Close	Volume
Date						
1981-05-01	0.126674	0.127790	0.126674	0.126674	0.097898	16553600
1981-05-04	0.126674	0.126674	0.126116	0.126116	0.097467	14448000
1981-05-05	0.126116	0.126116	0.125558	0.125558	0.097036	17539200
1981-05-06	0.122768	0.122768	0.122210	0.122210	0.094448	18950400
1981-05-07	0.123884	0.124442	0.123884	0.123884	0.095742	9363200
...
2024-04-23	165.350006	167.050003	164.919998	166.899994	166.899994	49537800
2024-04-24	166.539993	169.300003	166.210007	169.020004	169.020004	48251800
2024-04-25	169.529999	170.610001	168.149994	169.889999	169.889999	50558300
2024-04-26	169.880005	171.339996	169.179993	169.300003	169.300003	44525100
2024-04-29	173.369995	176.029999	173.100006	173.500000	173.500000	66162100

10839 rows × 6 columns

sp500_aapl.plot(grid=True,figsize=(8,5))

<Axes: xlabel='Date'>



```
sp500_aapl['Date'] = sp500_aapl.index
```

```
fig = px.area(  
    sp500_aapl,  
    x='Date',  
    y='Volume',  
    template='plotly_dark',  
    color_discrete_sequence=['cyan'],  
    title='AAPL index'  
)
```

```
fig.show()
```

The logarithmic return is a way of calculating the rate of return on an investment. To calculate it you need the initial value of the investment V_i , the final value V_f and the number of time periods t . You then take the natural logarithm of V_f divided by V_i , and divide the result by t :

$$R = \frac{\ln\left(\frac{V_f}{V_i}\right)}{t} \cdot 100\%$$

This value is normally expressed as a percentage, so you also multiply by 100.

The calculated rate will depend on the value of t that you use. If t is the number of years, then you get an annual rate. This then gives you the continuously compounded annual interest rate that you would need to receive in order to match the return on this investment.

Let p_t denote the price of an asset at time t . Then the return of an asset captures these relative movements and is defined as

$$r_t = \frac{p_t}{p_{t-1}} - 1$$

In words, a return is the change in price of an asset, relative to its previous value. Note that $p_t > 0$, and therefore $r_t > -1$.

In practice, “returns” often means “log returns”. Log returns are defined as

$$z_t = \log(1 + r_t) = \log\left(\frac{p_t}{p_{t-1}}\right) = \log(p_t) - \log(p_{t-1})$$

Log returns are normally distributed if prices are log normally distributed.

$$z \sim N(0, 1), \quad x = e^{\mu + \sigma z}$$

We could write this as

$$\log(x) = \mu + \sigma z \sim \mu + N(0, \sigma^2) \sim N(\mu, \sigma^2)$$

So we see that the prices are normally distributed. Then

$$z_t = \log(p_t) - \log(p_{t-1})$$

is the sum of two normal random variables.

```
sp500_aapl['Log Return'] = np.log(sp500_aapl['Adj Close']/sp500_aapl['Adj Close'].shift(1))  
  
print(sp500_aapl['Log Return'])
```

```
Date  
1981-05-01      NaN  
1981-05-04    -0.004415  
1981-05-05    -0.004434  
1981-05-06    -0.027027  
1981-05-07     0.013604  
...  
2024-04-23     0.006371  
2024-04-24     0.012622  
2024-04-25     0.005134  
2024-04-26    -0.003479  
2024-04-29     0.024505  
Name: Log Return, Length: 10839, dtype: float64
```


Histogram of log-returns

```
import seaborn as sns
```

```
plt.hist(sp500_aapl['Log Return'], color = 'blue', edgecolor = 'black',  
         bins = int(180/5))
```

```
plt.title('Histogram of Log>Returns')  
plt.xlabel('Log Return')  
plt.ylabel('Density')
```

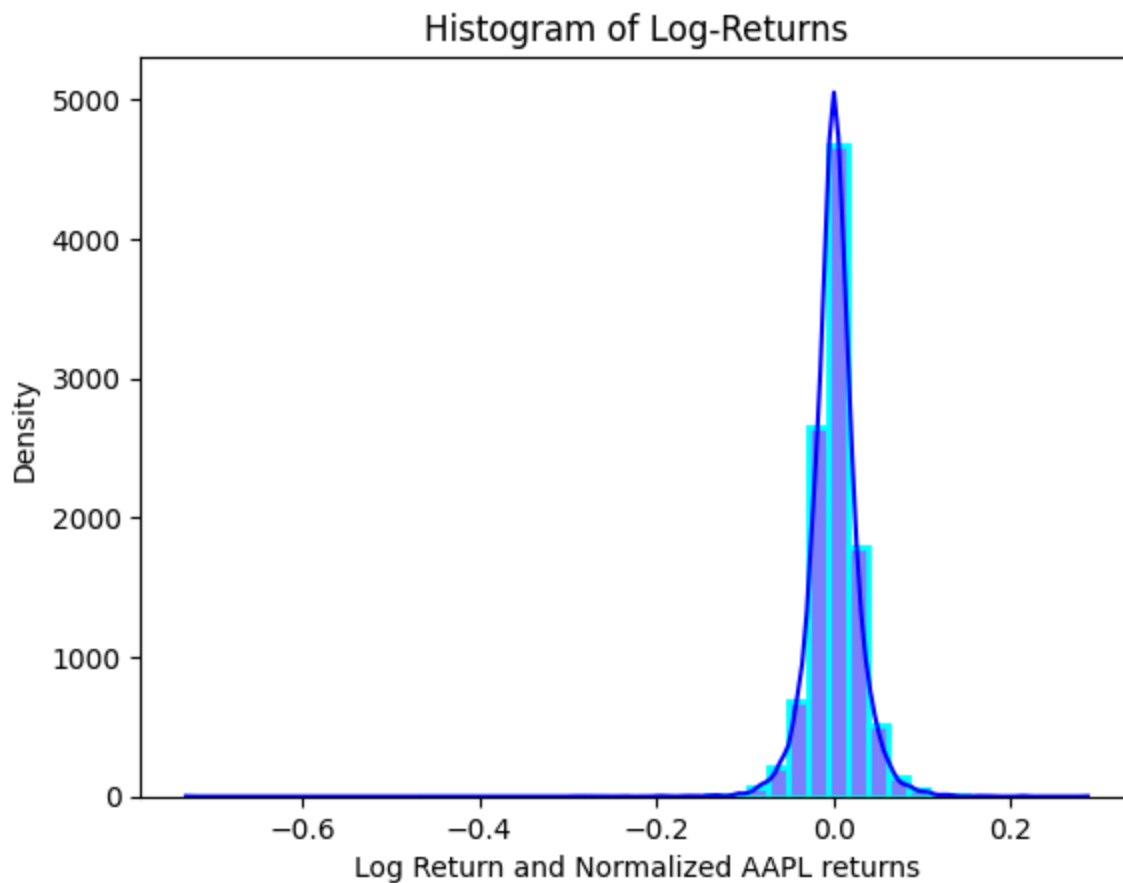
```
Text(0, 0.5, 'Density')
```



```
fig = px.histogram(  
    sp500_aapl,  
    x='Log Return',  
    color_discrete_sequence=['cyan'],  
    title='Histogram of Log-Returns',  
    labels={'Log Return': 'Log Return', 'Density': 'Density'},  
    template='plotly_dark'  
)  
  
fig.show()
```

We see that log-returns are normally distributed

```
sns.histplot(sp500_aapl['Log Return'], bins=int(180/4), color='blue',  
             edgecolor='cyan', kde=True, linewidth=2)  
  
# Set title and axis labels  
plt.title('Histogram of Log-Returns')  
plt.xlabel('Log Return and Normalized AAPL returns')  
plt.ylabel('Density')  
  
# Show the plot  
plt.show()
```



Calculating Daily Average Returns

Daily Average Returns are given by computing the mean of the log rate of return series.

```
daily_log_returns = sp500_aapl['Log Return'].mean()
daily_log_returns

0.0006901645243884324
```

Calculating Annual Average Returns

Annual Average Returns are given by computing the mean of the **log rate** of return series and then multiplying the value by 250 since 250 days exist in a business/trading day system. (because in one calendar year we have 250 workings days in a years; some people say that there are 252 tradings days, but we will consider 250 days, because it will convinient for us)

```
annual_log_returns = sp500_aapl['Log Return'].mean() * 250
annual_log_returns

0.17254113109710809

print(str(round(annual_log_returns*100, 2)) + ' %')

17.25 %
```



```
sp500_aapl['42d'] = np.round(sp500_aapl['Close'].rolling(window=42).mean(),2)
sp500_aapl['252d'] = np.round(sp500_aapl['Close'].rolling(window=252).mean(),2)
```

```
sp500_aapl.tail
```

```
<bound method NDFrame.tail of
Close
Date
1981-05-01    0.126674    0.127790    0.126674    0.126674    0.097898  \
1981-05-04    0.126674    0.126674    0.126116    0.126116    0.097467
1981-05-05    0.126116    0.126116    0.125558    0.125558    0.097036
1981-05-06    0.122768    0.122768    0.122210    0.122210    0.094448
1981-05-07    0.123884    0.124442    0.123884    0.123884    0.095742
...
2024-04-18  168.029999  168.639999  166.550003  167.039993  167.039993
2024-04-19  166.210007  166.399994  164.080002  165.000000  165.000000
2024-04-22  165.520004  167.259995  164.770004  165.839996  165.839996
2024-04-23  165.350006  167.050003  164.919998  166.899994  166.899994
2024-04-24  166.399994  169.289993  166.210007  168.574997  168.574997
```


```

Volume      Date  Log Return      42d      252d
Date
1981-05-01  16553600 1981-05-01         NaN      NaN      NaN
1981-05-04  14448000 1981-05-04    -0.004414      NaN      NaN
1981-05-05  17539200 1981-05-05    -0.004435      NaN      NaN
1981-05-06  18950400 1981-05-06    -0.027027      NaN      NaN
1981-05-07   9363200 1981-05-07     0.013605      NaN      NaN
...
2024-04-18  43122900 2024-04-18    -0.005731  173.62  181.18
2024-04-19  67772100 2024-04-19    -0.012288  173.23  181.17
2024-04-22  48116400 2024-04-22     0.005078  172.84  181.16
2024-04-23  48917700 2024-04-23     0.006371  172.42  181.17
2024-04-24  30881995 2024-04-24     0.009986  172.09  181.18
```

```
[10836 rows x 10 columns]>
```

```
sp500_aapl[['Close','42d','252d']].plot(grid=True,figsize=(8,5))
```

<Axes: xlabel='Date'>




```
sp500_aapl['Date'] = sp500_aapl.index
```

```
fig = px.line(  
    sp500_aapl,  
    x='Date',  
    y=['Close', '42d', '252d'],  
    template='plotly_dark',  
    color_discrete_sequence=['cyan', 'orange', 'green'],  
    title='Close Price with Moving Averages'  
)
```

```
fig.show()
```

```
sp500_aapl['Market Returns'] = np.log(sp500_aapl['Close'] / sp500_aapl['Close'].shift(1))  
sp500_aapl['Market Returns'].plot(grid=True,figsize=(8,5))
```

```
<Axes: xlabel='Date'>
```



```
sp500_aapl['Market Returns'] = np.log(sp500_aapl['Close'] / sp500_aapl['Close'].shift(1))
```

```
fig = px.line(x=sp500_aapl.index, y=sp500_aapl['Market Returns'], template='plotly_dark', color_discrete
```

```
fig.show()
```

```
sp500_aapl.keys()
```

```
Index(['Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume', 'Date',  
      'Log Return', '42d', '252d', 'Market Returns'],  
      dtype='object')
```

```
sp500_aapl['Date'] = sp500_aapl.index
```

```
fig = px.line(  
    sp500_aapl,  
    x='Date',  
    y='Close',  
    template='plotly_dark',  
    color_discrete_sequence=['cyan', 'orange', 'green'],  
    title='Close Price with Moving Averages'  
)
```

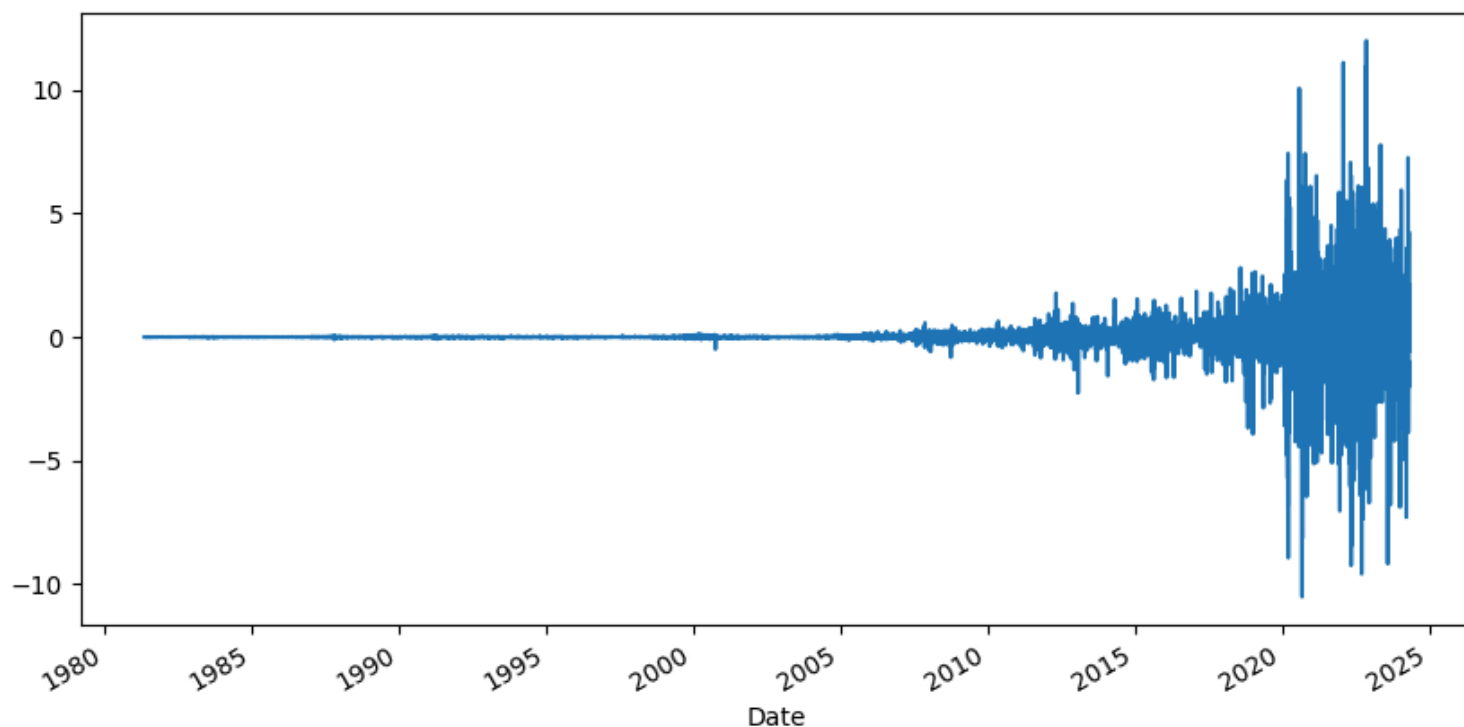
```
fig.show()
```

The first obvious thing to note, aside from the two giant dips at the tail end corresponding to the market crashes in 2020 and 2023, is that the data is clearly **non-stationary**.

Below we will try to calculate the first difference of the series, i.e we will subtract the previous value $t-1$ from the current value t to get the difference.

```
sp500_aapl['First Difference'] = sp500_aapl['Close'] - sp500_aapl['Close'].shift()
sp500_aapl['First Difference'].plot(figsize=(10, 5))
```

<Axes: xlabel='Date'>



```
fig = px.area(
    sp500_aapl,
    x='Date',
    y='First Difference',
    template='plotly_dark',
    color_discrete_sequence=['cyan'],
    title='First Difference'
)
```

```
fig.show()
```


The data no longer appears to be trending up over time and is instead centered around 0, but we have another **problem**.

If we pay attention to the variance., we will see that it's very small early on and steadily increases over time.

This is a sign that the data is not only non-stationary but also exponentially increasing. Now we'll apply a log transform to the original series.

```
sp500_aapl['Natural Log'] = sp500_aapl['Close'].apply(lambda x: np.log(x))
sp500_aapl['Natural Log'].plot(figsize=(10, 5))
```

<Axes: xlabel='Date'>



```
fig = px.area(
    sp500_aapl,
    x='Date',
    y='Natural Log',
    template='plotly_dark',
    color_discrete_sequence=['cyan'],
    title='Natural Log'
)
```

```
fig.show()
```

```
sp500_aapl['Original Variance'] = pd.rolling_var(sp500_aapl['Close'], 30, min_periods=None, freq=None)
sp500_aapl['Log Variance'] = pd.rolling_var(sp500_aapl['Natural Log'], 30, min_periods=None, freq=None)
```

```
fig, ax = plt.subplots(2, 1, figsize=(13, 12))
sp500_aapl['Original Variance'].plot(ax=ax[0], title='Original Variance')
sp500_aapl['Log Variance'].plot(ax=ax[1], title='Log Variance')
fig.tight_layout()
```

```

-----
AttributeError                                Traceback (most recent call last)
Cell In[139], line 1
----> 1 sp500_aapl['Original Variance'] = pd.rolling_var(sp500_aapl['Close'], 30,
min_periods=None, freq=None, center=True, how = None)
      2 sp500_aapl['Log Variance'] = pd.rolling_var(sp500_aapl['Natural Log'], 30,
min_periods=None, freq=None, center=True, how = None)
      4 fig, ax = plt.subplots(2, 1, figsize=(13, 12))

AttributeError: module 'pandas' has no attribute 'rolling var'

```

Чтобы изменить содержимое ячейки, дважды нажмите на нее (или выберите "Ввод")


Now we'll calculate the first difference from the logged series, i.e we'll calculate log-returns

```

sp500_aapl['Logged First Difference'] = sp500_aapl['Natural Log'] - sp500_aapl['Natural Log'].shift()
sp500_aapl['Logged First Difference'].plot(figsize=(10, 5))

```

<Axes: xlabel='Date'>



A stationary sequence is a random sequence whose joint probability distribution is invariant over time. If a random sequence X_j is stationary then the following holds:

$$F_{X_0, X_1, \dots, X_m}(x_0, x_1, \dots, x_m) = F_{X_k, X_{k+1}, \dots, X_{k+m}}(x_0, x_1, \dots, x_m),$$

where F is the joint cumulative distribution function of the random variables in the subscript, i.e

$$(X_0, X_1, \dots, X_m) \stackrel{d}{\sim} (X_k, X_{k+1}, \dots, X_{k+m})$$

If a sequence is stationary then it has a constant mean (which may not be finite): $E(X[n]) = \mu$ for all n

\newline \newline


Now we have a stationary time series model of daily changes to the S&P 500 index Apple's company (this is just an assumed hypothesis, where next we'll try to show it).

Below we'll create some lag variables $y(t-1)$, $y(t-2)$ etc. and examine their relationship to $y(t)$ and look at 1 and 2-day lags along with weekly and monthly lags to look for "seasonal" effects.

```
sp500_aapl['Lag 1'] = sp500_aapl['Logged First Difference'].shift()
sp500_aapl['Lag 2'] = sp500_aapl['Logged First Difference'].shift(2)
sp500_aapl['Lag 5'] = sp500_aapl['Logged First Difference'].shift(5)
sp500_aapl['Lag 30'] = sp500_aapl['Logged First Difference'].shift(30)
```

One interesting visual way to evaluate the relationship between lagged variables is to do a scatter plot of the original variable vs. the lagged variable and see where the distribution lies with a joint plot using the seaborn package.

```
sns.jointplot(data=sp500_aapl, x='Logged First Difference', y='Lag 1', palette='Set2', kind = 'reg')
plt.show()
```



Notice how tightly packed the mass is around 0, besides it also appears to be pretty evenly distributed - the marginal distributions on both axes are roughly normal.

This seems to indicate that knowing the index value one day doesn't tell us much about what it will do the next day.

```
from statsmodels.tsa.stattools import acf
from statsmodels.tsa.stattools import pacf

lag_correlations = acf(sp500_aapl['Logged First Difference'].iloc[1:])
lag_partial_correlations = pacf(sp500_aapl['Logged First Difference'].iloc[1:])
```


The auto-correlation function computes the correlation between a variable and itself at each lag step up to some limit (in this case 40).

The partial auto-correlation function computes the correlation at each lag step that is NOT already explained by previous, lower-order lag steps.

We can plot the results to see if there are any significant correlations.


```
fig, ax = plt.subplots(figsize=(10, 6))
ax.plot(lag_correlations, marker='o', linestyle='--')
```

[<matplotlib.lines.Line2D at 0x1fe101d7dd0>]




```
fig, ax = plt.subplots(figsize=(10, 6))
ax.plot(lag_partial_correlations, marker='o', linestyle='--')
```

[<matplotlib.lines.Line2D at 0x1fe1012d450>]



The autocorrelation and partial-autocorrelation results are very close to each other (the first upper plot is the autocorrelation results and the second lower plot is the partial-autocorrelation results).

What this shows is that there is no significant correlation between the value at time t and at any time prior to t up to 40 steps behind. (the correlation coefficient of 0.2 before excluding outliers is considered as negligible correlation while 0.3 after excluding outliers may be interpreted as weak positive correlation).

In other words, the log-returns series is a random walk.

!!! Now I have little issues because I needed to show that log-returns are/aren't stationary time series, but I showed that log-returns series is a random walk !!!

Is there any statement or proof that we'll show that any random walk is non-stationary / random walk is stationary time series?

In the world website web I've read one fact that a random walk is a non-stationary, but I have no clue how to prove it mathematically correctly.

I saw the proof of this fact in this link: <https://spureconomics.com/random-walk-model-and-stationarity/>, but I am not sure about it.

If it's true, i.e a log-returns which is a random walk cannot be a stationary time series.

```
%matplotlib inline
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import statsmodels.api as sm
import seaborn as sb
sb.set_style('darkgrid')

path = os.getcwd() + '\data\stock_data.csv'
stock_data = pd.read_csv(path)
stock_data['Date'] = stock_data['Date'].convert_objects(convert_dates='coerce')
stock_data = stock_data.sort_index(by='Date')
stock_data = stock_data.set_index('Date')
```

FileNotFoundError

Traceback (most recent call last)

Cell In[96], line 11

```
      8 sb.set_style('darkgrid')
     10 path = os.getcwd() + '\data\stock_data.csv'
--> 11 stock_data = pd.read_csv(path)
     12 stock_data['Date'] = stock_data['Date'].convert_objects(convert_dates='coerce')
     13 stock_data = stock_data.sort_index(by='Date')
```

File

```
~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-
packages\Python311\site-packages\pandas\io\parsers\readers.py:912, in
read_csv(filepath_or_buffer, sep, delimiter, header, names, index_col, usecols, dtype,
engine, converters, true_values, false_values, skipinitialspace, skiprows, skipfooter, nrows,
na_values, keep_default_na, na_filter, verbose, skip_blank_lines, parse_dates,
infer_datetime_format, keep_date_col, date_parser, date_format, dayfirst, cache_dates,
iterator, chunksize, compression, thousands, decimal, lineterminator, quotechar, quoting,
doublequote, escapechar, comment, encoding, encoding_errors, dialect, on_bad_lines,
delim_whitespace, low_memory, memory_map, float_precision, storage_options, dtype_backend)
     899 kwds_defaults = _refine_defaults_read(
     900     dialect,
     901     delimiter,
     (... )
     908     dtype_backend=dtype_backend,
     909 )
     910 kwds.update(kwds_defaults)
--> 912 return _read(filepath_or_buffer, kwds)
```